

Paper overview

CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training

Jianmin Bao, Dong Chen, Fang Wen, Houqiang Li, Gang Hua

ICCV 2017, Venice, Italy DOI: 10.1109/ICCV.2017.299

- ▶ **Problem:** conditional generators struggle to produce images that are realistic, diverse, and faithful to a fine-grained class.
- ▶ **Motivation:** VAE outputs tend to blur; GAN training can suffer from unstable gradients and mode collapse.
- ▶ **Main contribution:** a conditional VAE-GAN with asymmetric objectives, mean feature matching, and pairwise reconstruction constraints.
- ▶ **Project code:** <https://github.com/sivajeetsabdakar/cvae-gan-pytorch-reproduction> **Pending**





Figure 1. **Synthesized images** using our CVAE-GAN model at high resolution (128×128) for different classes. The generated samples are realistic and diverse within a class.

Paper Figure 1: 128×128 class-conditioned samples.

Methodology



Paper Figure 3: encoder E , generator G , classifier C , discriminator D .

- **Workflow:** $E(x, c)$ infers latent z ; $G(z, c)$ reconstructs or samples x' ; C preserves class identity; D enforces realism.
- **Datasets:** FaceScrub, Oxford 102 Flowers, and CUB-200 birds.
- **Generator objective:** combine pixel reconstruction, discriminator-feature matching, and classifier-feature matching.
- **Asymmetric training:** C and D use cross-entropy, while G matches feature means with an ℓ_2 objective.
- **Key settings:** 128×128 images, 256-D latent code, batch normalization, $\lambda_1 = 3$, $\lambda_2 = 1$, $\lambda_3 = \lambda_4 = 10^{-3}$.

Reproduction and implementation

Algorithm 1 The training pipeline of the proposed CVAE-GAN algorithm.

Require: m , the batch size. n , class number. θ_E , initial E network parameters. θ_G , initial G network parameters. θ_D , initial D network parameters. θ_C , initial C network parameters, $\lambda_1 = 3$, $\lambda_2 = 1$, $\lambda_3 = 10^{-3}$ and $\lambda_4 = 10^{-3}$.

```
1: while  $\theta_G$  has not converged do
2:   Sample  $\{x_r, c_r\} \sim P_r$  a batch from the real data;
3:    $\mathcal{L}_C \leftarrow -\log(P(c_r|x_r))$ 
4:    $z \leftarrow E(x_r, c_r)$ 
5:    $\mathcal{L}_{KL} \leftarrow KL(q(z|x_r, c_r)||P_z)$ 
6:    $x_f \leftarrow G(z, c_r)$ 
7:   Sample  $z_p \sim P_z$  a batch of random noise, sample  $c_p$  a
   batch of random classes;
8:    $x_p \leftarrow G(z_p, c_p)$ 
9:    $\mathcal{L}_D \leftarrow -(\log(D(x_r)) + \log(1 - D(x_f)) + \log(1 - D(x_p)))$ 
10:  Calculate  $x_r$  feature center  $\frac{1}{m} \sum_i^m f_D(x_r)$  and  $x_p$  feature
   center  $\frac{1}{m} \sum_i^m f_D(x_p)$ ;
11:   $\mathcal{L}_{GD} \leftarrow \frac{1}{2} \|\frac{1}{m} \sum_i^m f_D(x_r) - \frac{1}{m} \sum_i^m f_D(x_p)\|_2^2$ 
12:  Calculate each class  $c_i$  feature center  $f_C^{c_i}(x_r)$  for  $x_r$  and
    $f_C^{c_i}(x_p)$  for  $x_p$  using moving average method;
13:   $\mathcal{L}_{GC} \leftarrow \frac{1}{2} \sum_{c_i} \|f_C^{c_i}(x_r) - f_C^{c_i}(x_p)\|_2^2$ 
14:   $\mathcal{L}_G \leftarrow \frac{1}{2} (\|x_r - x_f\|_2^2 + \|f_D(x_r) - f_D(x_f)\|_2^2 + \|f_C(x_r) -
   f_C(x_f)\|_2^2)$ 
15:   $\theta_C \leftarrow \theta_C + \nabla_{\theta_C}(\mathcal{L}_C)$ 
16:   $\theta_D \leftarrow \theta_D + \nabla_{\theta_D}(\mathcal{L}_D)$ 
17:   $\theta_G \leftarrow \theta_G + \nabla_{\theta_G}(\lambda_2 \mathcal{L}_G + \lambda_3 \mathcal{L}_{GD} + \lambda_4 \mathcal{L}_{GC})$ 
18:   $\theta_E \leftarrow \theta_E + \nabla_{\theta_E}(\lambda_1 \mathcal{L}_{KL} + \lambda_2 \mathcal{L}_G)$ 
19: end while
```

Paper Algorithm 1: alternating updates for C , D , G , and E .

- **Original toolchain:** Torch; GoogLeNet encoder, six-layer deconvolutional generator, DCGAN discriminator, AlexNet classifier.
- **Dataset preparation:** align FaceScrub faces from five landmarks; crop Oxford Flowers using masks; retain original CUB-200 images; resize inputs to 128×128 .
- **Implementation approach:** alternate network updates exactly as Algorithm 1 and use moving averages for class-specific feature centers.
- **Experimental setup:** train all networks from scratch; sample $z \sim \mathcal{N}(0, I)$ for category-conditioned testing.
- **Workspace status:** no source code, checkpoints, or logs were present in LAB03. Framework changes and executed settings are therefore **Pending**.

Results and comparison



Paper Figure 5: real samples, CVAE, CGAN, FM-CGAN, and CVAE-GAN. CVAE-GAN retains category identity and within-class diversity more consistently than the baselines shown.

Original quantitative results

Method	Top-1 acc.	Realism
Real data	99.61%	20.85
CVAE	8.09%	10.29
CGAN	61.97%	15.79
FM-CGAN	79.76%	19.40
CVAE-GAN	97.78%	19.03

Original vs. reproduced

Original evaluation: 53k generated faces.

Reproduced metrics: Pending

Difference: not computable without a run.

Metrics reproduced from paper Table 1. "Realism" is the paper's Inception Score variant.

Conclusion and learning

- ▶ **Reproduction status:** the paper analysis and presentation are complete; model reproduction remains **Pending** because the workspace contains no implementation or run artifacts.
- ▶ **Major finding:** feature matching stabilizes conditional GAN training, while the encoder and pairwise losses protect diversity and class structure.
- ▶ **Reported impact:** synthetic identity augmentation improved LFW verification from 91.87% to 92.98% in the paper.
- ▶ **Challenges:** legacy Torch code, three datasets with different preprocessing, four-network optimization, and evaluation with a separate classifier.
- ▶ **Next improvements:** implement in PyTorch, record seeds and checkpoints, compare FID/KID alongside the legacy score, and run ablations for each loss term.
- ▶ **Takeaway:** CVAE-GAN ties each latent code to a real example and trains the generator with feature statistics instead of the discriminator's cross-entropy alone.




Figure 8. Visualization comparison between different generator G , each trained with different combination of losses.

Method	Training Data	Accuracy
no data augmentation	80K	91.87%
existing identities augmentation	80K + 100K	92.77%
5k new identities augmentation	80K + 500K	92.98%

Table 2. Results of face data augmentation.

Paper Figures 7–8 and Table 2: inpainting, loss ablations, and face augmentation.